Responsible AI Model Card

Model: Logistic Regression on the UCI Adult Income dataset.
Purpose: Educational income classification and fairness audit.
Baseline accuracy: 0.842
Baseline fairness metrics: {'Disparate Impact': np.float64(0.3300683371298405), 'Statistical Parity Difference': np.float64(-0.17199073746563104), 'Equal Opportunity Difference': np.float64(-0.09047658334020847), 'Average Odds Difference': np.float64(-0.081714392888473 5)}
After reweighing: {'Disparate Impact': np.float64(0.5764907276296798), 'Statistical Parity Difference': np.float64(-0.0922056800910864), 'Equal Opportunity Difference': np.float64(0.11752743173005531), 'Average Odds Difference': np.float64(0.04962329778915463)}
Limitations: Historical US data may encode social bias. Removing gender does not remove proxy bias. This model is not intended for real-world eligibility or employment decisions.